




# Using Fork and Pipe

ECE 650  
Methods & Tools for Software Engineering (MTSE)  
Fall 2023

Presented by  
Dr. Albert Wasef



Used by permission from Prof. Arie Gurfinkel



UNIVERSITY OF  
**WATERLOO**

# Additional Information


## Advanced Linux Programming

- Chapter 2.1 (Interacting with Execution Environment)
- Chapter 3 (Processes)
- Chapter 5.4 (Pipes)

The book is available from the links below

<https://github.com/MentorEmbedded/advancedlinuxprogramming/blob/gh-pages/alp-folder/advanced-linux-programming.pdf>

<https://github.com/MentorEmbedded/advancedlinuxprogramming/tree/gh-pages>

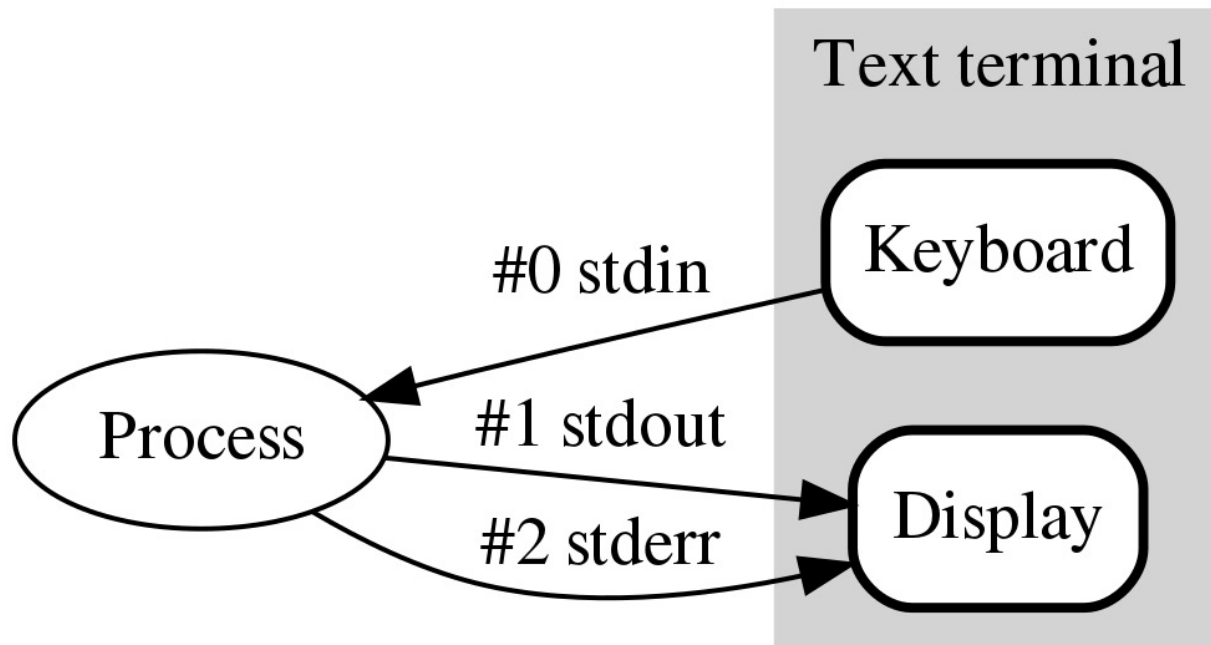


# PROCESS

# What is a “Process”?

- What is a process:
  - *“A running instance of a program”*
  - *Examples:*
    - *Each of the two instances of Chrome*
    - *The shell and the ls command executed, each is a process*
- Advanced programmers use multiple processes to
  - Do several tasks at once
  - Increase robustness (one process fails, other still running)
  - Make use of already-existing processes

# Standard input, output, and error



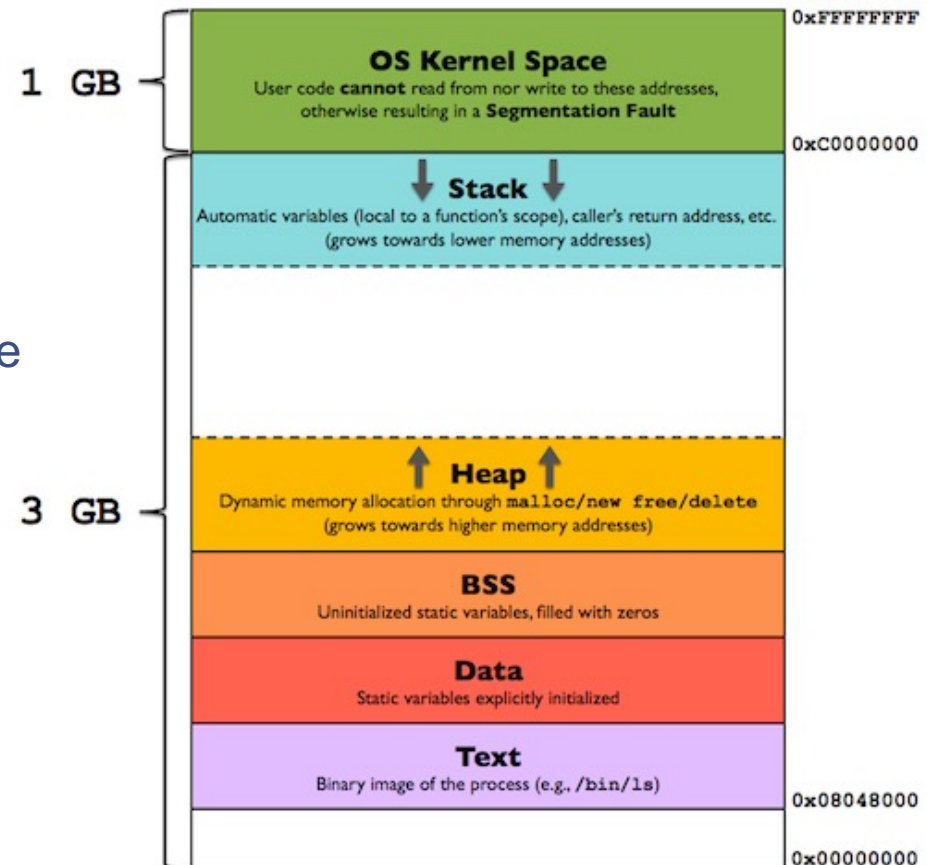
<https://en.wikipedia.org/wiki/File:Stdstreams-notitle.svg>

- Let's change stdin, stdout, and stderr

# The “Guts” of a Process!


Process of a 32bit architecture

- The main components of a process:
  - An executable piece of code (a program)
  - Data that is input or output by the program
  - Execution context (information about the program needed by OS)



<https://gabrieletolomei.wordpress.com/miscellanea/operating-systems/in-memory-layout/>

# Life Cycle of a (Unix) Process



Idle

state when the process is creating

Ready

ready to run

Run

executing

Wait

waiting for resources (CPU, disk, network, etc.)

Zombie

ended, waiting to be collected

# Let's Dissect a Process!

- Windows:
  - Task manager
- Unix-like (Mac and Linux):
  - In the terminal type:
    - `ps` or `top` or `htop`
    - `ps -f` for full details



# System Calls: using OS provided services

System calls (a.k.a., **syscalls**) is the programmatic way for a program (in user-space) to request services from an operating system (in kernel-space)

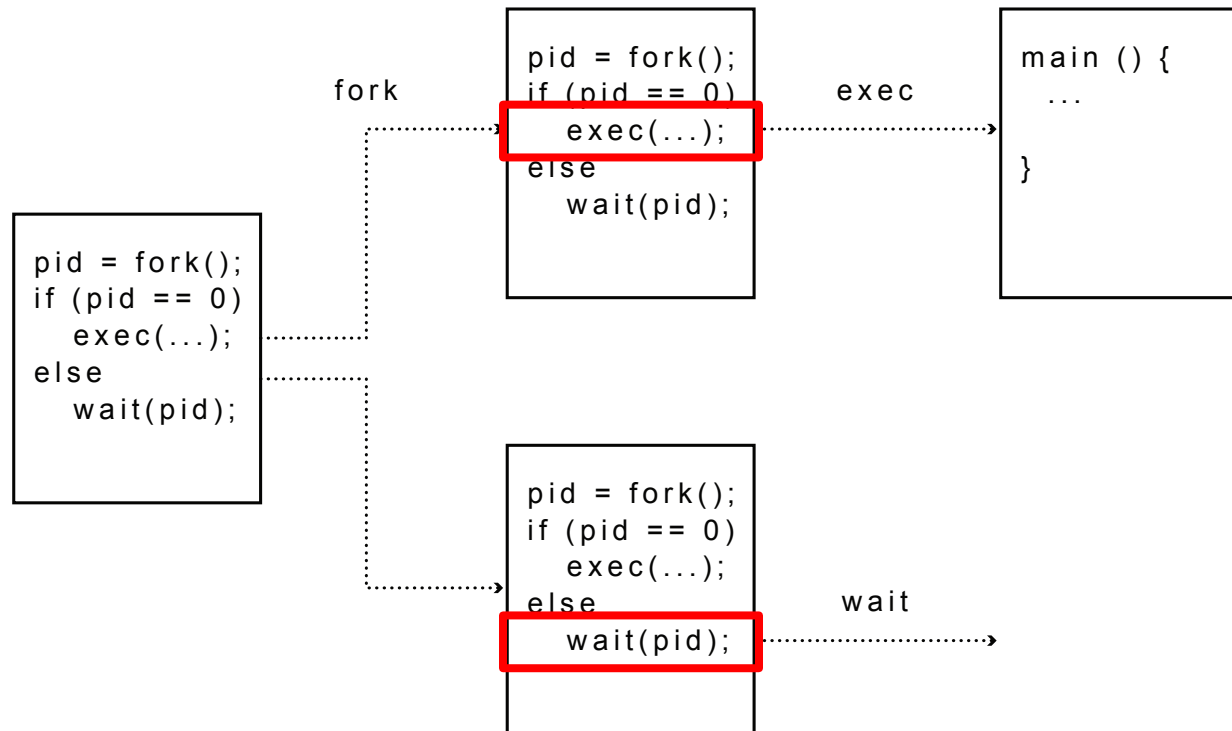
In code, a syscall looks like a function calls

- however, unlike a regular call it transfer control to OS
- therefore, syscalls are **independent** of the programming language
  - e.g., syscall `fork()` works the **same** in Python, C, C++, ...
- but **depend** on the OS
  - e.g., `fork()` works **differently** on Windows, Linux, OSX

Syscalls are used to access OS managed resources

- processes, files, memory, devices, ...

# UNIX Process Management



# UNIX Process Management System Calls

## fork()

- Create a copy of current process and start it as a child

## execv() / execl() / ...

- Load an executable into the current process and run it

## wait() / waitpid()

- Wait for a child process to finish

## kill()

- Send a signal (e.g., **SIGTERM**, **SIGKILL**, **SIGINT**) to another process

# The Parent of a Process

- Each process (with some exceptions) has a parent process (indicated by `ppid` – parent process identifier)
- Can we get this information within a program?
  - YES!
  - Use `getpid()` and `getppid()` libc functions defined in `unistd.h`


```
// c++ -o main main.cpp
#include <unistd.h>
#include <iostream>

int main(void) {
    std::cout << "my pid: " << getpid()
               << " ppid: " << getppid() << "\n";
}
```

# KILLing a Process!

- Run `kill` in the terminal (run kill with -KILL)

## HANDLING NON-RESPONDING & FROZEN APPLICATIONS



[https://www.reddit.com/r/linuxmasterrace/comments/3y42qz/killing\\_a\\_non\\_responding\\_process/](https://www.reddit.com/r/linuxmasterrace/comments/3y42qz/killing_a_non_responding_process/)

# Creating a Process - `fork()` system call

Forks an execution of the process

- after a call to `fork()`, a new process is created (called child)
- the original process (called parent) continues to execute concurrently
- in the parent, `fork()` returns the process `id` of the child that was created
- in the child, `fork()` return `0` to indicate that this is a child process
- The parent and child are independent

Man(ual) Page

- `man 2 fork`

# exec() – executing a program in a process

`exec()` series of functions are used to start another program in the current process

- after a call to `exec()` the current process is replaced with the image of the specified program
- different versions allow for different ways to pass command line arguments and environment settings
- `int execl(const char *file, char *const argv[ ])`
  - `file` is a path to an executable
  - `argv` is an array of arguments. By convention, `argv[0]` is the name of the program being executed

Man page

- `man 3 exec`

# kill() – sending a signal

A process can send a signal to any other process

- usually the parent process sends signals to its children
- `int kill(pid_t pid, int sig)`
  - send a signal `sig` to a process `pid`
- useful signal: `SIGTERM`
  - asks a process to terminate

When a parent process exits, the children processes are terminated

It's a good practice to kill and wait for children to terminate before exiting

Man page

- `man 2 kill`



# Signals

- A special message sent to a process
- Signals are asynchronous
- Different types of signals (defined in `signal.h`)
  - `SIGTERM`: Termination
  - `SIGINT`: Terminal interrupt (Ctrl+C)
  - `SIGKILL`: Kill (can't be caught or ignored)
  - `SIGBUS`: BUS error
  - `SIGSEGV`: Invalid memory segment access
  - `SIGPIPE`: Write on a pipe with no reader, Broken pipe
  - `SIGSTOP`: Stop executing (can't be caught or ignored)
- Handling a signal:
  - Default *disposition*
  - Signal handler procedure
- Sending signal from one process to another process (`SIGTERM`, `SIGKILL`)

# waitpid() – Waiting for a child

A parent process can wait for a child process to terminate

- `pid_t waitpid(pid_t pid, int *status, int options)`
  - block until the process with the specified `pid` terminates
  - the return code from the terminating process is placed in `status`
  - `options` control whether the function blocks or not
    - 0 is a good choice for options

Man page

- `man 2 waitpid`

# pipe() and dup2() – Inter-Process Communication

`pipe()` creates a ONE directional pipe

- two file descriptors: one to write to and one to read from the pipe
- a process can use the pipe by itself, but this is unusual
- typically, a parent process creates a pipe and shares it with a child, or between multiple children
- some processes read from it, and some write to it
  - there can be multiple writers and multiple readers
    - although multiple writers is more common

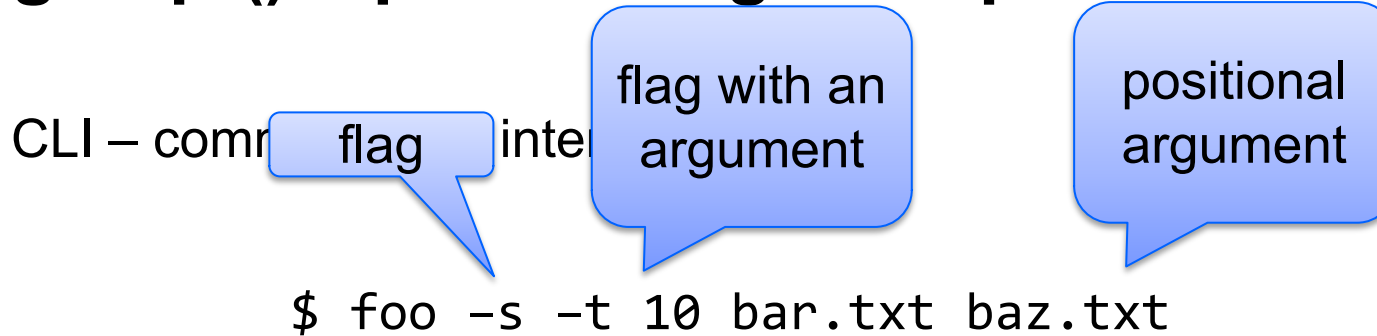
`dup2()` duplicates a file descriptor

- used to redirect standard input, standard output, and standard error to a pipe (or another file)
- `STDOUT_FILENO` is the number of the standard output

Man pages

- `man 2 pipe`
- `man 2 dup2`

# getopt() – processing CLI options



At a start of the program, `main(argc, argv)` is called, where

- `argc` is the number of CLI arguments
- `argv` is an array of 0 terminated strings for arguments
  - e.g., `argv[0]` is “foo”, `argv[1]` is “-s”, `argv[2]` is “-t”, `argv[3]` is “10”, ...

`getopt()` is a library function to parse CLI arguments

- `getopt(argc, argv, “st:”)`
- input: arguments and a string describing desired format
- output: returns the next argument and an option value
- see example in `using_getopt.cpp`

# /dev/urandom – Really Random Numbers

`/dev/urandom` is a special file (device) that provides supply of “truly” random numbers

“infinite size file” – every read returns a new random value

To get a random value, read a byte/word from the file

see `using_rand.cpp` for an example

**Have to use it for Assignment 3!**

